

# An experimental study of the learnability of congestion control

Anirudh Sivaraman, Keith Winstein, Pratiksha Thaker,  
Hari Balakrishnan

MIT CSAIL

May 25, 2014

# Designing congestion-control protocols today

- ▶ Formulate a mental model of the target network and application workload
- ▶ Decide on the protocol's goal
- ▶ Design a protocol to achieve this goal on the target network
- ▶ Can either be implicit or explicit

# But, the model is always wrong!

- ▶ Lost throughput due to stochastic loss
- ▶ Bufferbloat when queues are incorrectly sized
- ▶ Incast in datacenters

# Our work

- ▶ Can we formalize this design process?
- ▶ Quantify the consequences of model mismatch?

# Contributions

- ▶ Formalize learnability in the context of congestion control.
- ▶ Use it to answer:
  - ▶ Do we need to know the link speed exactly?
  - ▶ What is the cost of backwards compatibility?
  - ▶ Do we need to know the topology exactly?

# Approach

- ▶ Specify a *training scenario*.
  - ▶ Topology
  - ▶ Locations of senders and receiver
  - ▶ Application workload
  - ▶ Buffer size and queuing discipline
- ▶ Specify an *objective function*.
- ▶ Synthesize protocol automatically.
- ▶ Evaluate on a *testing scenario* inside ns-2

# Automated protocol synthesis

- ▶ Find best protocol, given an imperfect network model
- ▶ The problem is hard to solve in general
- ▶ Rely on Remy [?] to produce congestion-control protocols.

# Caveats and non-goals

- ▶ Very simple, controlled experiments
- ▶ Results could change with better protocol-design tools
- ▶ Not trying to understand Remy's internals



# Is using Remy reasonable?

- ▶ Training scenario:

Link speed	32 Mbits/sec
------------	--------------

Minimum RTT	150 ms
-------------	--------

Topology	Dumbbell
----------	----------

Number of senders	2
-------------------	---

Workload	1 sec ON/OFF times
----------	--------------------

Buffer size	5 BDP
-------------	-------


Objective function	$\sum \log(\text{throughput}) - \log(\text{delay})$
--------------------	---

- ▶ Testing scenario identical to training scenario


# Is using Remy reasonable?

- ▶ A hypothetical centralized scheme (CEN)
  - ▶ Every time a sender turns ON/OFF, solve  $\sum \log(\text{throughput})$
  - ▶ Set each sender's rate using obtained solution
  - ▶ Zero queuing delay


# Is using Remy reasonable?



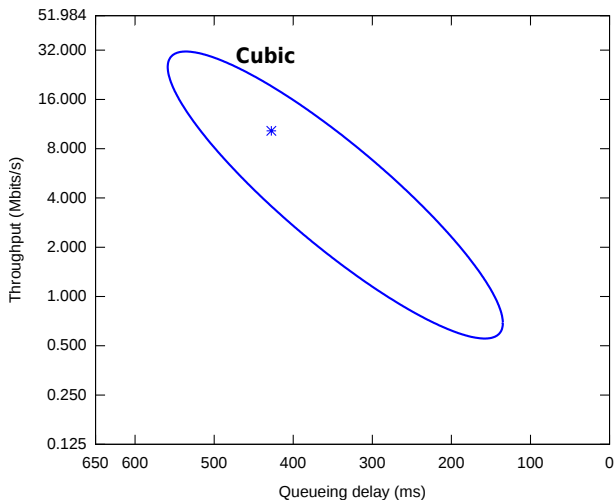
# Is using Remy reasonable?




# Is using Remy reasonable?



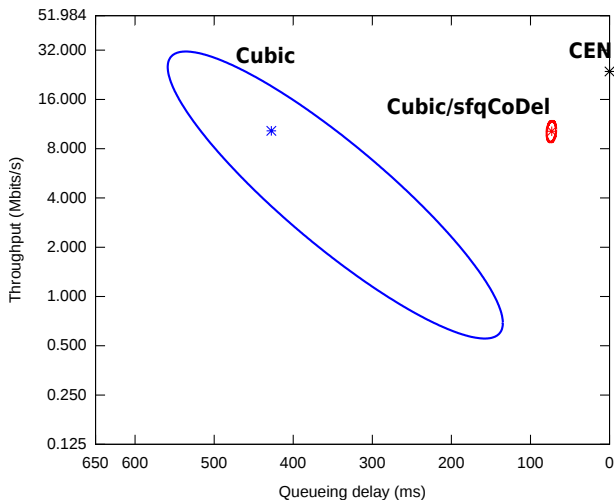
# Is using Remy reasonable?



# Is using Remy reasonable?

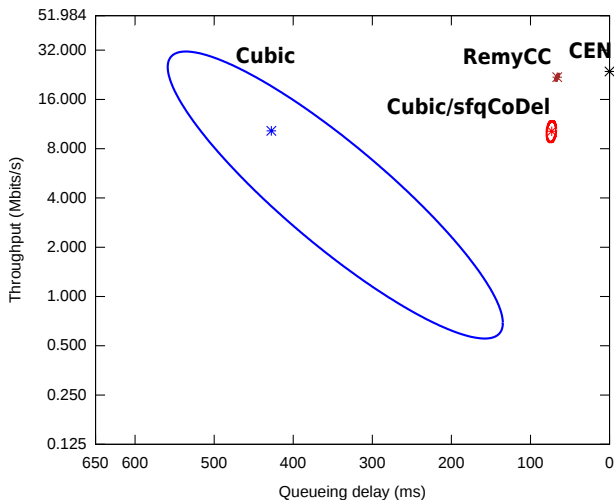


# Is using Remy reasonable?





# Is using Remy reasonable?



# The cost of generality, or forwards-compatibility


<b>RemyCC</b>	<b>Link rates</b>	<b>RTT</b>	<b>Senders</b>	<b>ON/OFF time</b>	<b>Topology</b>
1000x	1–1000 Mbps	150 ms	2	1 sec	Dumbbell
100x	3.2–320 Mbps	150 ms	2	1 sec	Dumbbell
10x	10–100 Mbps	150 ms	2	1 sec	Dumbbell
2x	22–44 Mbps	150 ms	2	1 sec	Dumbbell

**Table :** Training scenarios for forwards-compatibility experiment


<b>Link rates</b>	<b>RTT</b>	<b>Senders</b>	<b>ON/OFF time</b>	<b>Topology</b>
1–1000 Mbps	150 ms	2	1 sec	Dumbbell

**Table :** Testing scenario for forwards-compatibility experiment


# The cost of generality, or forwards-compatibility




# The cost of generality, or forwards-compatibility




# The cost of generality, or forwards-compatibility




# The cost of generality, or forwards-compatibility



# The cost of generality, or forwards-compatibility




# The cost of generality, or forwards-compatibility







# The cost of generality, or forwards-compatibility




# The cost of generality, or forwards-compatibility



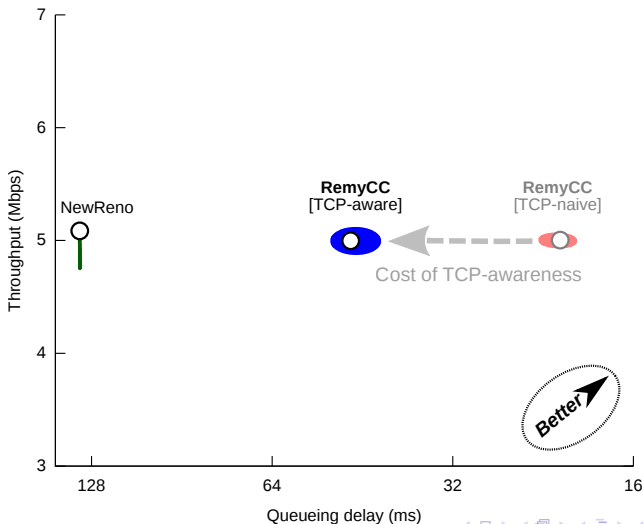
# RemyCC competing against itself




# RemyCC competing against itself




# RemyCC competing against itself



# RemyCC competing against TCP NewReno



# RemyCC competing against TCP NewReno









# When the model is wrong about the topology

One bottleneck




# When the model is wrong about the topology

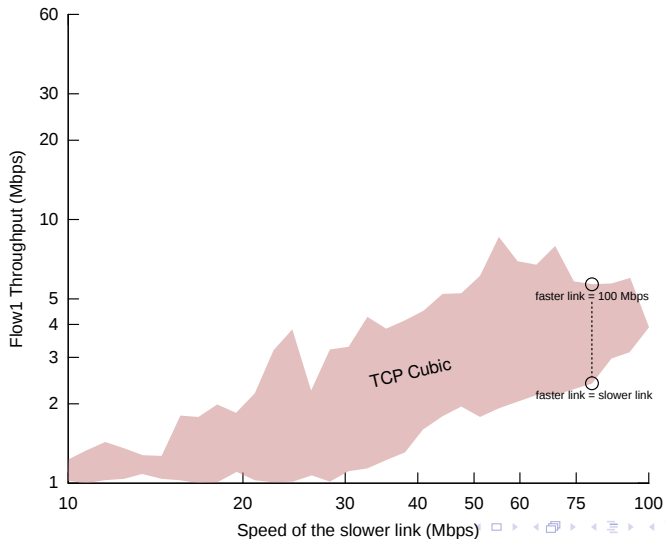
Two bottlenecks




# When the model is wrong about the topology




# When the model is wrong about the topology



# When the model is wrong about the topology



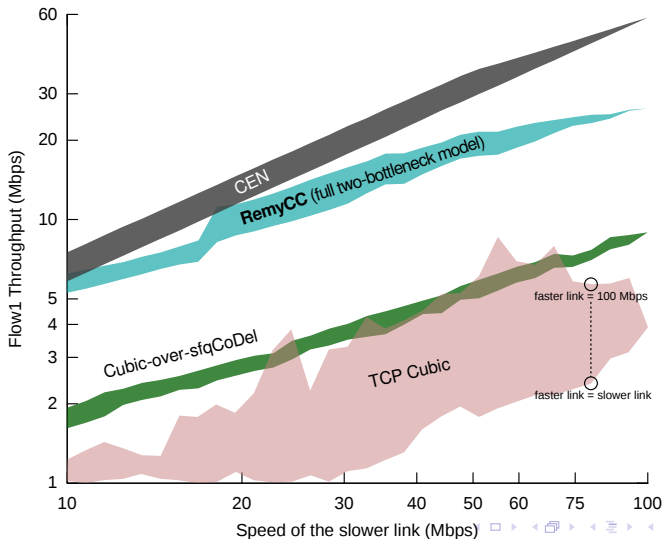
# When the model is wrong about the topology




Anirudh Sivaraman, Keith Winstein, Pratiksha Thaker, Hari Balakrishnan

An experimental study of the learnability of congestion control

# When the model is wrong about the topology



# When the model is wrong about the topology





# Related Work

- ▶ Probably approximately correct learning
- ▶ Transfer learning
- ▶ Machine-generated congestion control

# Limitations and future work

- ▶ Generalizability to more complex topologies?
- ▶ Better characterization of gap from optimal
- ▶ Do results change if we learn in-network behavior as well?
- ▶ Model mismatches between simulation and the real world

# Backup slides

# Can applications with different objectives coexist?

- ▶ Tpt. Sender: A throughput-intensive sender


$$\log(\text{throughput}) - 0.1 * \log(\text{delay}) \quad (1)$$

- ▶ Lat. Sender: A latency-sensitive sender


$$\log(\text{throughput}) - 10.0 * \log(\text{delay}) \quad (2)$$

- ▶ Running over a FIFO queue


# Training for diversity has a cost ...




# Training for diversity has a cost ...



# Training for diversity has a cost ...




but, benefits the docile sender





but, benefits the docile sender



but, benefits the docile sender

