

# An experimental study of the learnability of congestion control

Anirudh Sivaraman, Keith Winstein, Pratiksha Thaker,  
Hari Balakrishnan

MIT CSAIL

May 23, 2014

# Designing congestion-control protocols today

- ▶ Formulate a mental model of the target network and application workload
- ▶ Decide on the protocol's goal
- ▶ Design a protocol to achieve this goal on the target network
- ▶ Can either be implicit or explicit

# But, the model is always wrong!

- ▶ Lost throughput due to stochastic loss
- ▶ Bufferbloat when queues are incorrectly sized
- ▶ Diminished fairness in small-packet regimes
- ▶ Incast in datacenters

# Our work

- ▶ Can we formalize this design process?
- ▶ Quantify the consequences of model mismatch?

# Approach

- ▶ Specify a *training scenario*.
  - ▶ Topology
  - ▶ Locations of senders and receiver
  - ▶ Application workload
  - ▶ Buffer size and queuing discipline
- ▶ Specify an *objective function*.
- ▶ Synthesize protocol automatically.
- ▶ Evaluate on a *testing scenario* inside ns-2

# Automated protocol synthesis

- ▶ Find best protocol, given an imperfect network model.
- ▶ Unfortunately, problem is NEXP-complete.

# Tractable Attempts at Optimal

- ▶ Rely on Remy [?] to produce Tractable Attempts at Optimal (TAO) congestion-control protocols.
- ▶ Approaches upper bounds on throughput and lower bounds on delay.

# How far off is Remy from the optimal?

- ▶ Training scenario:

Link speed	32 Mbits/sec
------------	--------------

minimum RTT	150 ms
-------------	--------

Topology	Dumbbell
----------	----------

Number of senders	2
-------------------	---

Workload	1 sec ON/OFF times
----------	--------------------


Buffer size	5 BDP
-------------	-------

Objective function	$\sum \log(\text{throughput}) - \log(\text{delay})$
--------------------	---

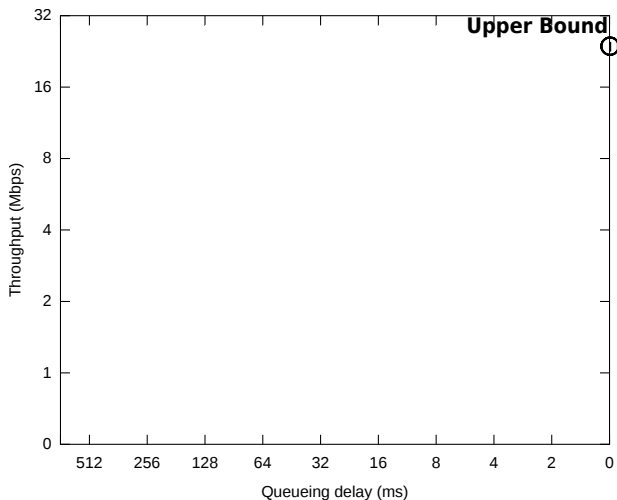
- ▶ Testing scenario identical to training scenario



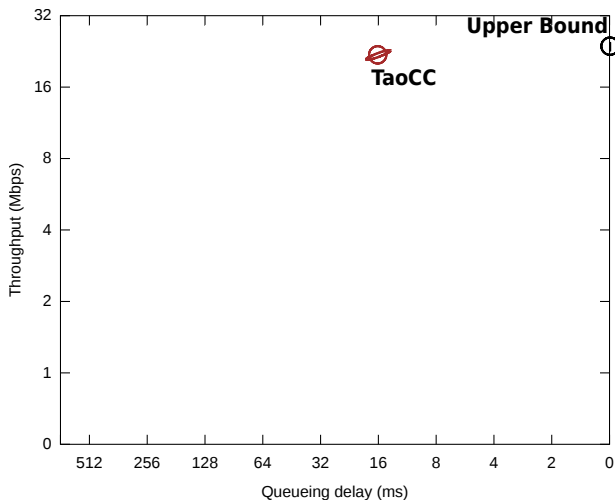
# How far off is Remy from the optimal?



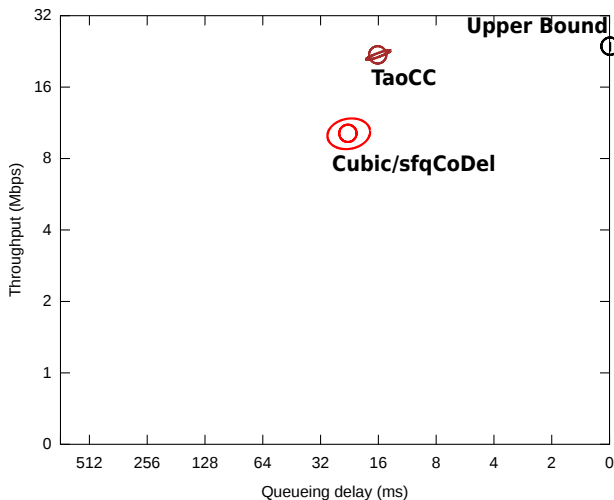
# How far off is Remy from the optimal?




# How far off is Remy from the optimal?



# How far off is Remy from the optimal?



# How far off is Remy from the optimal?




# The cost of generality, or forwards-compatibility


<b>Tao</b>	<b>Link rates</b>	<b>RTT</b>	<b>Senders</b>	<b>ON/OFF time</b>	<b>Topology</b>
1000x	1–1000 Mbps	150 ms	2	1 sec	Dumbbell
100x	3.2–320 Mbps	150 ms	2	1 sec	Dumbbell
10x	10–100 Mbps	150 ms	2	1 sec	Dumbbell
2x	22–44 Mbps	150 ms	2	1 sec	Dumbbell

Table : Training scenarios for forwards-compatibility experiment

# The cost of generality, or forwards-compatibility




# The cost of generality, or forwards-compatibility







# The cost of generality, or forwards-compatibility




# The cost of generality, or forwards-compatibility




# The cost of generality, or forwards-compatibility



# The cost of generality, or forwards-compatibility




# The cost of generality, or forwards-compatibility




# When the model is wrong about the topology

One bottleneck




# When the model is wrong about the topology

Two bottlenecks

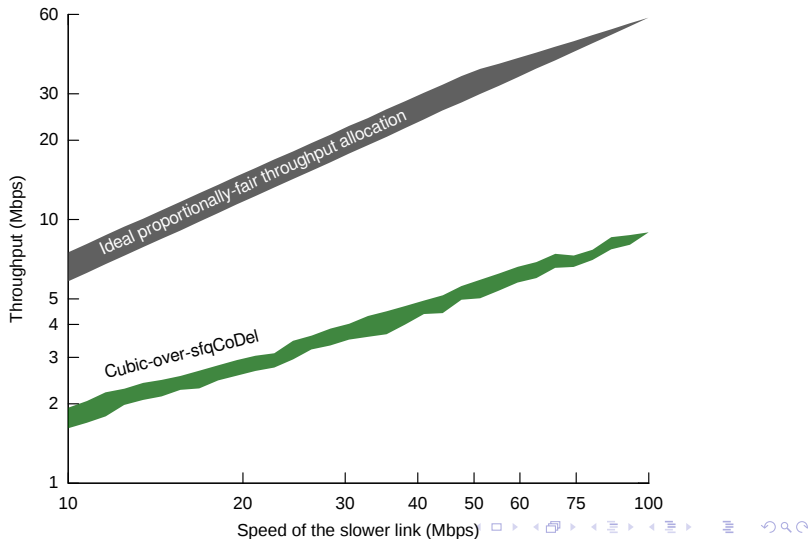


# When the model is wrong about the topology







# When the model is wrong about the topology




# When the model is wrong about the topology




# When the model is wrong about the topology




# RemyCC competing against itself




# RemyCC competing against itself




# RemyCC competing against itself



# RemyCC competing against TCP NewReno



# RemyCC competing against TCP NewReno







# Can applications with different objectives coexist?

- ▶ Tpt. Sender: A throughput-intensive sender


$$\log(\text{throughput}) - 0.1 * \log(\text{delay}) \quad (1)$$

- ▶ Lat. Sender: A latency-sensitive sender

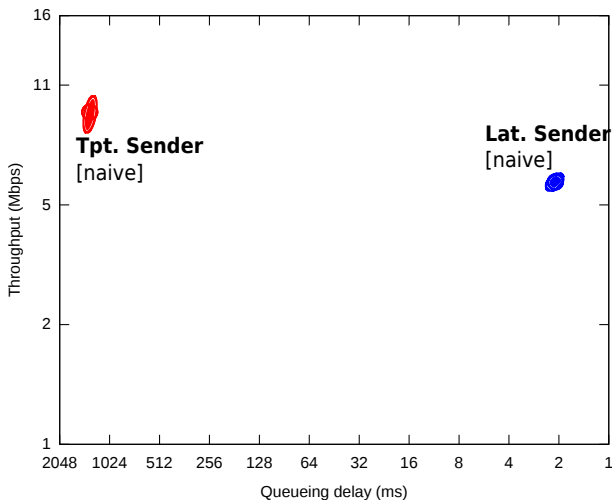
$$\log(\text{throughput}) - 10.0 * \log(\text{delay}) \quad (2)$$

- ▶ Running over a FIFO queue


# Training for diversity has a cost ...




# Training for diversity has a cost ...




# Training for diversity has a cost ...



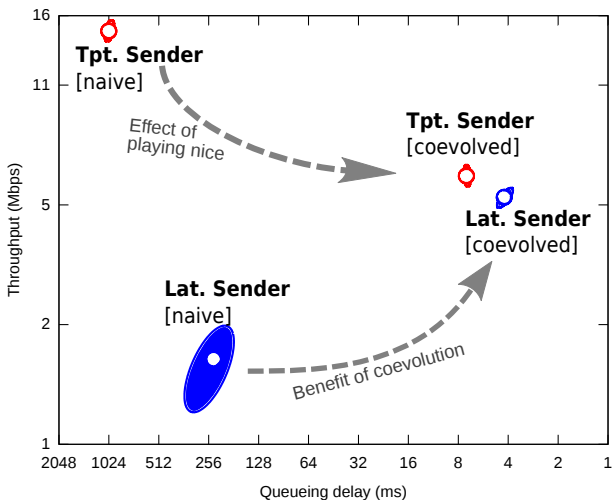
but, benefits the docile sender



but, benefits the docile sender



but, benefits the docile sender





# Related Work

- ▶ Probably approximately correct learning
- ▶ Transfer learning
- ▶ Machine-generated congestion control

# Limitations and future work

- ▶ Generalizability to more complex topologies?
- ▶ Better characterization of gap from optimal
- ▶ Do results change if we learn in-network behavior as well?
- ▶ Model mismatches between simulation and the real world